

On Solving the Multiple Variable Gapped Longest Common Subsequence Problem

Marko Djukanović^{1,2,6}

Aleksandar Kartelj⁴

Nikola Balaban²

Sašo Džeroski⁵

Christian Blum³

Žiga Zebec⁶

¹University of Nova Gorica, Nova Gorica, Slovenia

²Faculty of Natural Sciences and Mathematics, University of Banja Luka, Banja Luka, B&H

³Artificial Intelligence Research Institute (IIIA-CSIC), Barcelona, Spain

⁴Faculty of Mathematics, University of Belgrade, Belgrade, Serbia

⁵Jozef Stefan Institute, Ljubljana, Slovenia

⁶Institute of Information Sciences (IZUM), Maribor, Slovenia

- EUROCAST 2026: 20th International Conference on Computer Aided Systems Theory, February 23-27, 2026, Las Palmas de Gran Canaria, Spain –



Sofinancira
Evropska unija



Outline

- Introduction
- Problem Definition
- Rooted graph state space
- Iterative multi-source Beam Search
- Experimental Evaluation: General & special problem
- Conclusions

Introduction

- Objects we deal with: sequences (strings) over finite alphabet Σ
 - DNA/RNA over {A, T, G, C/U}
 - Proteins over 20 (canonical) amino acids: {A, C, D, E, P, Q...}
- **One of central tasks in computational biology:**
 - sequence comparison, finding common motifs between sequences
 - compare structurally but also semantically/functionality
 - sequence alignment problems
- **Subsequences:** reveal structural similarities → **Longest common subsequence problem** variants (studied 50 years already)

Longest common subsequence problem (LCSP)

Definition (LCSP)

Input: Given an arbitrary set of sequences $S = \{s_1, \dots, s_m\}$

Task: Find a subsequence s **common** for all sequences from S with **maximum possible length** ($|s|$).

Example

Input: $S = \{\text{AATTGC, ATTAC}\}$

LCS solution: $s = \text{ATTC}$

- Basic problem in computational biology, well-solved theoretically and practically

LCS: Literature & Problem Variants

- When $m = 2$: polynomially solvable (in $O(n^2)$):
 - Dynamic programming, Hunt-Szymanski, Hirschenberg, ...
- When m arbitrary large – \mathcal{NP} -hard:
 - subject of interest within last 30 years – approximation approaches, meta-heuristics (ACO, Beam search, ...), but also exact approaches (A*, anytime approaches, DAG-based, ...)

Problem-related practical variants:

- Arc-annotated LCS problem
- Constrained, Repetition-free, filled LCS problem, ...
- Gapped LCS problem

The gapped LCS problem

Definition (A gap sequence)

Given is a sequence s and an assigned function $G_s: \{1, \dots, |s|\} \mapsto \mathbb{N}$. A pair (s, G_s) is called a gap sequence.

Definition (A gapped subsequence)

Sequence \tilde{s} is a **gapped subsequence** of (s, G_s) iff

- \tilde{s} is a subsequence of s
- the gap constraint G_s is fulfilled w.r.t. *positions of embedding letters* of \tilde{s} in s
 - suppose $i_1, \dots, i_{|\tilde{s}|}$ are those **positions**
 - $(\forall j = 2, \dots, |\tilde{s}|) i_j - i_{j-1} \leq G_s(i_j) + 1$ (consecutive positions close enough)

Problem definition

Example

$s = \text{AATTGC}$, $G_s(\cdot) = 1$

- $\tilde{s} = \text{ATG}$, the embedding: AATTGC (**valid** gapped subsequence)
- $\tilde{s} = \text{ATG}$, the embedding: AATTGC (**invalid** gapped subsequence)

Definition (The multiple (variable) gapped LCS problem – MVGLCSP)

Input: Given is a set of gapped sequences $\{(s_1, G_{s_1}), \dots, (s_m, G_{s_m})\}$.

Task: Find the longest subsequence \tilde{s} so that

- \tilde{s} is common subsequence of each s_i
- \tilde{s} fulfills all gap constraints G_{s_i} ($i = 1, \dots, m$)

Note: when $G_{s_i} = n$ (the length of longest sequence) \Rightarrow VGLCSP = LCSP.

Literature & Motivation for VGLCSP

- Peng and Yang (2012, 2014): studied the $m = 2$ (poly) version by **three dynamic programming** (DP) approaches (basic one, two advanced)
- Manea et al. (2024): Complexity bounds, (parameterised) complexity analysis investigated
- **NP-hard** under arbitrary large m

Motivation:

- **Genetics and molecular biology**: applications in DNA/protein analysis where variable structural distances between residues must be respected
- **Time-series analysis**: in settings where events are required to occur within specified temporal delays (Lainscsek et al. (2015))

Methodology

- Based on the significant extension of the state space graph formulation for LCS problem (Djukanovic et al. (2020))
- **Gap constraints:** incorporated to cut-off invalid extensions (edges) among the LCS extensions immediately
- Many root/source nodes in the state graph (the position of leading letter in a common subsequence free to choose)

Root-based state graph formulation: rough idea

- Each **state** $v = (p^L, l^v)$: one or more feasible partial solutions where
 - a vector of positions p^L refer to the positions of suffixes of input sequences relevant to further expand these sols
 - the length of current partial solution l^v
- **Expansion** of v : extend (concatenate) partial solutions feasibly (and non-dominantly) by one letter in all possible ways, respecting gap constraints
- **Non-expandable nodes**: complete solutions;

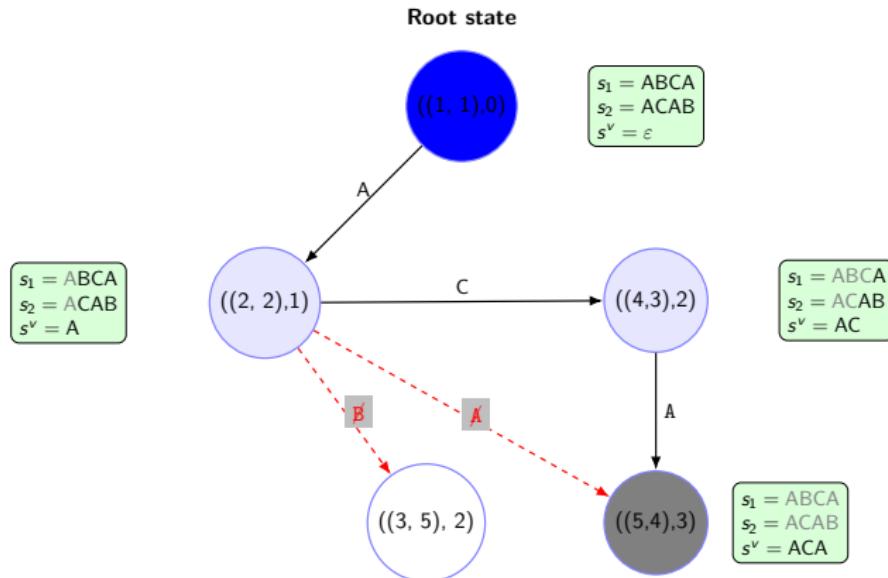
Decision: select appropriate (match) p^L for start (empty solution)

⇒ possibly (**exponentially**) many root nodes

$\text{Space}(p^L)$: root-based state (sub) graph induced by root node $(p^L, 0)$.

Root-based state graph formulation: example

$\text{Space}(r = ((1, 1), 0))$ for MVGLCSP for $s_1 = \boxed{A}BCA$, $s_2 = \boxed{A}CAB$, assuming $G_1 = G_2 = 1$:



An issue with the root-state-space formulation: example

Example

$S = \{s_1 = \text{ATGG}[\text{A}]AA, s_2 = \text{ATCC}[\text{A}]AA\}$, with gap constraints
 $G_{s_1} = G_{s_2} = 1$. In this instance, any state with position vector $\mathbf{p}^L = (5, 5)$ cannot be reached from the initial state $((1, 1), 0)$ by standard direct transitions.


Consequence: $((5, 5), 0) \notin \text{Space}((1, 1)) \Rightarrow$ The optimal common subsequence AAA is unreachable

How to fix it?

Iterative multi-source beam search (IMSBS) strategy

- Explicitly enumerating all root states computationally prohibitive ($O(n^m)$ time)
- **IMSBS** (rough idea):
 - Exploring Space(r): **beam search** — limited breadth-first-search (BFS) strategy; parameter β controls the number of nodes to be further pursued
 - Dynamically explore multiple promising regions (Space(\cdot)) of the state space
 - **Iteratively identify** a set of **promising** candidate root states

Workflow of the IMSBS



Heuristic guidances in BS

Three LCS heuristic guidances used:

- “Look-ahead” for the remaining sequence length: $\text{UB}_1(v)$
- *Character Frequency Alignment* score: UB_2
 - sum up the maximum possible occurrences of each letters in subsequences
- *A probability-based heuristic* guidance: h_{prob} (the pre-processed matrices of probabilities (Mousavi and Tabataba (2012)))

Experimental Evaluation: arbitrary large m -case

- Bs: a baseline beam search approach, allowing only a single iteration of IMSBS (utilizing a huge β)
- IMSBS-GREEDY: fix beam-width $\beta = 1$ for the forward BS of IMSBS, perform a large number of beam search (impact of iterations on the overall performance of IMSBS)
- IMSBS: a tuned version; configured to use an average runtime comparable to that of Bs

Benchmark set RANDOM

For each combination of instance parameters

- $n \in \{50, 100, 200, 500\}$
- $m \in \{2, 3, 5, 10\}$
- $|\Sigma| \in \{2, 4\}$

10 random problem instances are generated (sequences uniformly at random).

The gap constraints generated from
 $G_s(\cdot) \in \mathcal{U}(\{ \lfloor 0.5 \cdot |\Sigma| \rfloor, \dots, \lfloor 1.5 \cdot |\Sigma| \rfloor \})$.

⇒ A total of **320 problem instances** is generated.

Parameter tuning of IMSBS

We fixed (less sensible) params:

- Bs (backward): $\beta' = 10$, and UB₂ (efficient)
- Candidate root nodes from \mathcal{R} ordered by UB₂ (decreasingly)
- At each iteration, **10 best nodes** taken from \mathcal{R} as the initial beam

Tuned parameters:

- β (in BS-forward)
- Heuristic guidance in BS (forward)
 - {UB₁, UB₂, h_{prob} }

Grid search used to tune; avg. solution quality over all (320) instances!

The results of tuning

- **(Baseline)** Bs $\Rightarrow \beta = 10,000$ and $h = h_{\text{prob}}$
- **Imsbs** $\Rightarrow h = \text{UB}_2$ and $\beta = 500$, and $\text{beam_iter} = 100$.
- **Imsbs-Greedy:** $\beta = 1$, $\text{beam_iter} = 10,000$ (comparable or a bit larger avg. runtime to that of IMSBS)

Numerical results

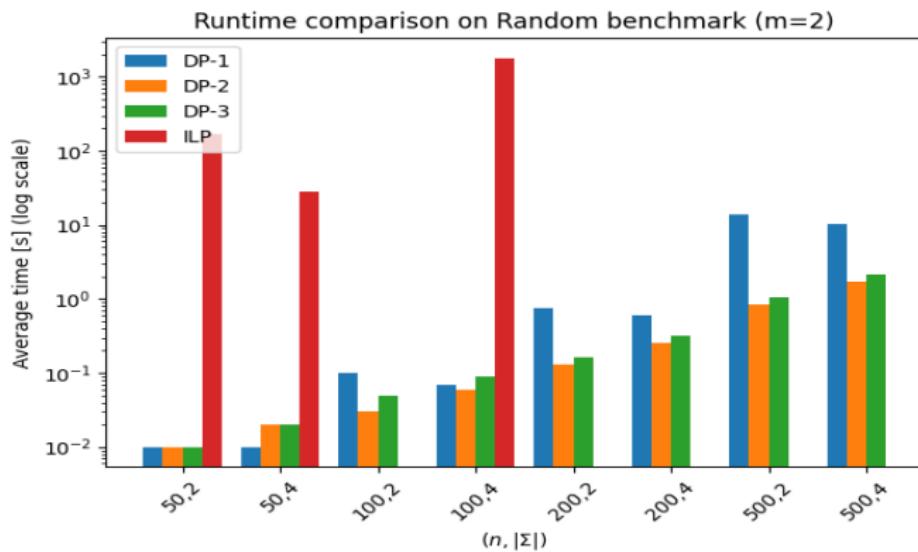
Inst.			Bs		IMSBs-GREEDY		IMSBs	
<i>m</i>	<i>n</i>	$ \Sigma $	\overline{obj}	$\overline{t}[s]$	\overline{obj}	$\overline{t}[s]$	\overline{obj}	$\overline{t}[s]$
2	50	2	33.6	0.02	33.1	0.00	37.7	0.06
2	50	4	30.1	0.98	27.7	0.00	30.1	0.16
2	100	2	48.9	2.07	64.5	0.01	72.8	0.94
2	100	4	62.1	11.19	56.9	0.01	61.6	0.91
2	200	2	99.1	18.56	95.5	0.02	136.4	6.21
2	200	4	120.5	38.15	116.1	0.05	124.9	6.58
2	500	2	65.3	23.27	153.6	0.07	265.7	119.75
2	500	4	214.6	163.69	227.7	0.12	294.4	60.49
3	50	2	17.5	0.03	27.2	0.00	31.2	0.14
3	50	4	21.7	0.18	21.5	0.00	22.9	0.19
3	100	2	19.7	0.06	41.8	0.01	58.5	3.15
3	100	4	34.1	2.35	43.4	0.03	48.4	5.58
3	200	2	15.2	0.16	63.6	0.02	90.0	22.48
3	200	4	85.3	23.45	77.2	0.08	97.1	72.25
3	500	2	12.6	0.10	69.9	0.07	102.9	53.56
3	500	4	90.7	86.35	104.2	0.29	187.7	412.27
5	50	2	4.8	0.00	14.9	0.00	20.0	0.36
5	50	4	8.9	0.01	13.6	0.06	15.3	0.79
5	100	2	6.3	0.01	17.7	0.01	22.4	0.57
5	100	4	5.3	0.01	23.2	10.85	22.1	1.44
5	200	2	5.3	0.01	21.6	0.03	26.6	1.07
5	200	4	6.4	0.02	32.5	604.11	25.7	2.10
5	500	2	5.9	0.10	25.5	0.14	27.9	3.22
5	500	4	6.8	0.10	43.6	1341.25	26.9	3.52
10	50	2	1.7	0.00	8.8	2.28	9.1	0.47
10	50	4	1.9	0.00	7.0	508.36	6.1	1.46
10	100	2	1.1	0.00	14.0	1421.10	8.6	0.54
10	100	4	2.2	0.01	8.9	1800.45	6.3	1.51
10	200	2	2.5	0.01	13.2	1710.49	10.3	0.77
10	200	4	2.2	0.02	7.9	1800.54	6.1	1.74
10	500	2	1.8	0.08	13.8	1611.70	9.5	1.53
10	500	4	1.9	0.09	8.2	1800.46	6.1	2.37
Avg.			32.38	10.97	46.82	394.14	59.73	24.63

Numerical results for $m = 2$ case


- DP-1: the basic DP ($O(n^2m^2)$)
- DP-2: an advanced DP, uses Incremental Suffix Maximum Queries (ISMQ) — Col and All matrices ($O(n^2 + mn)$)
- DP-3: an enhanced DP — ISMQ handled with a *deque* structure (slightly modified w.r.t. the literature)
- ILP: an integer linear programming, **proposed in this work**, motivated by the ILP model for LCSP

The first known empirical comparison for the $m = 2$ (80 instances).

Numerical evaluation: $m = 2$



The $m = 2$ case: heuristic performance vs. optimal solution



Relative solution quality achieved by the heuristic approaches compared to the optimal solutions (shown in %).

Conclusions and Future Work

- Proposed a **general heuristic framework** IMSBS to solve the multiple VGLCS problem
- Combines Beam search calls (backward-and-forward manner) in an **iterative way** while producing promising source nodes for further BS iterations
 - Balancing intensification and diversification
- Empirical studies conducted for the first time on the **synthetic instances**: IMSBS wins over the baseline Beam search

Future work:

- **Real-world** instance-case scenario
- Lack of more advanced heuristic guidance: **Data-driven/ML heuristic** involving various local and global features (NN-based)

Thank you for your attention!